

What's TensorFlow™?

- Open source software library for numerical computation using data flow graphs
- Originally developed by Google Brain Team to conduct machine learning and deep neural networks research
- General enough to be applicable in a wide variety of other domains as well

TensorFlow provides an extensive suite of functions and classes that allow users to build various models from scratch.

Why TensorFlow?

- Python API
- Portability: deploy computation to one or more CPUs or GPUs in a desktop, server, or mobile device with a single API
- Flexibility: from Raspberry Pi, Android, Windows, iOS, Linux to server farms
- Visualization (TensorBoard is da bomb)
- Checkpoints (for managing experiments)
- Auto-differentiation *autodiff* (no more taking derivatives by hand. Yay)
- Large community (> 10,000 commits and > 3000 TF-related repos in 1 year)
- Awesome projects already using TensorFlow

Companies using Tensorflow

- Google
- OpenAI
- DeepMind
- Snapchat
- Uber
- Airbus
- eBay
- Dropbox
- A bunch of startups


```
import tensorflow as tf
```



Graphs and Sessions

Data Flow Graphs


TensorFlow separates definition of computations from their execution



Data Flow Graphs

Phase 1: assemble a graph

Phase 2: use a session to execute operations in the graph.



What's a tensor?

What's a tensor?

An n-dimensional array

0-d tensor: scalar (number)

1-d tensor: vector

2-d tensor: matrix

and so on

Data Flow Graphs

Visualized by TensorBoard

```
import tensorflow as tf
```


```
a = tf.add(3, 5)
```

Why x, y?

TF automatically names the nodes when you don't explicitly name them.

```
x = 3
```

```
y = 5
```



Data Flow Graphs


Interpreted?

```
import tensorflow as tf
```

```
a = tf.add(3, 5)
```

Nodes: operators, variables, and constants

Edges: tensors



Data Flow Graphs

Interpreted?

```
import tensorflow as tf
```


```
a = tf.add(3, 5)
```

Nodes: operators, variables, and constants

Edges: tensors

Tensors are data.

Data Flow -> Tensor Flow (I know, mind=blown)



Data Flow Graphs


```
import tensorflow as tf
```

```
a = tf.add(3, 5)
```

```
print a
```

```
>> Tensor("Add:0", shape=(), dtype=int32)
```

(Not 8)



How to get the value of a?

Create a **session**, assign it to variable sess so we can call it later

Within the session, evaluate the graph to fetch the value of a

How to get the value of a?

Create a **session**, assign it to variable `sess` so we can call it later

Within the session, evaluate the graph to fetch the value of `a`


```
import tensorflow as tf
```

```
a = tf.add(3, 5)
```

```
sess = tf.Session()
```

```
print sess.run(a) tf.print(a)
```

```
sess.close()
```



The session will look at the graph, trying to think: hmm, how can I get the value of `a`, then it computes all the nodes that leads to `a`.

How to get the value of a?

Create a **session**, assign it to variable `sess` so we can call it later

Within the session, evaluate the graph to fetch the value of `a`


```
import tensorflow as tf
```

```
a = tf.add(3, 5)
```

```
sess = tf.Session()
```

```
print sess.run(a)      >> 8
```

```
sess.close()
```



The session will look at the graph, trying to think: hmm, how can I get the value of `a`, then it computes all the nodes that leads to `a`.

How to get the value of a?

Create a **session**, assign it to variable `sess` so we can call it later

Within the session, evaluate the graph to fetch the value of `a`

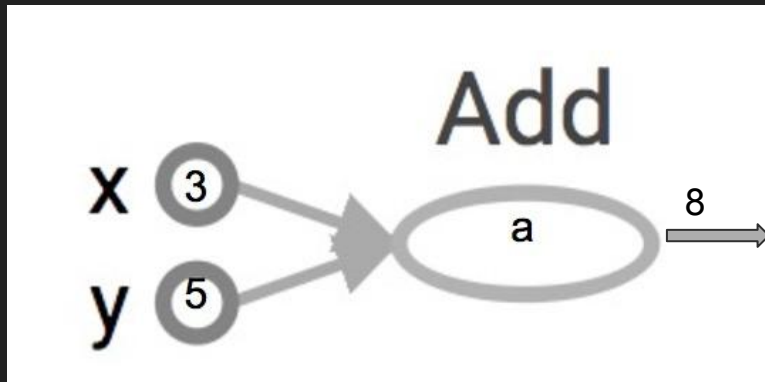
```
import tensorflow as tf

a = tf.add(3, 5)

sess = tf.Session()

with tf.Session() as sess:
    print sess.run(a)

sess.close()
```




tf.Session()

A Session object encapsulates the environment in which Operation objects are executed, and Tensor objects are evaluated.

More graphs

Visualized by TensorBoard


```
x = 2  
y = 3  
  
op1 = tf.add(x, y)  
op2 = tf.mul(x, y)  
op3 = tf.pow(op2, op1)  
  
with tf.Session() as sess:  
    op3 = sess.run(op3)
```



Subgraphs

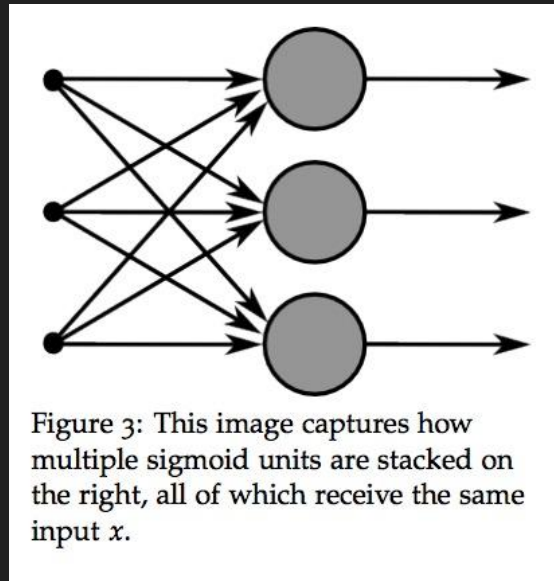
Possible to break graphs into several chunks and run them parallelly across multiple CPUs, GPUs, or devices

Example: AlexNet



Why graphs

1. Save computation (only run subgraphs that lead to the values you want to fetch)
2. Break computation into small, differential pieces to facilitates auto-differentiation
3. Facilitate distributed computation, spread the work across multiple CPUs, GPUs, or devices
4. Many common machine learning models are commonly taught and visualized as directed graphs already



A neural net graph by Richard Socher (CS224D)